

Practical Machine Learning Course Project

Liu Yi Hung

2015/9/24

Introduction

This paper is aim to predict personal activity, the data is from devices such as *Jawbone Up*, *Nike FuelBand*, and *Fitbit*. More informations about data can be found at <http://groupware.les.inf.puc-rio.br/har>. The data is downloaded from here, [training data](#) and [testing data](#), soucre: <http://groupware.les.inf.puc-rio.br/har>.

In this project, I'll use variables to predict the outcome, which is the *classe* variable in the data. The algorithm I used is *random forest*, more informations about random forest can be found at [Random forests](#). Basically, the main analysis take the advantage of the *caret* and *randomForest* packages in R.

Why I choose this algorithm is because the advantage of random forest, these features is copied from (Random Forests: Leo Breiman and Adele Cutler)[https://www.stat.berkeley.edu/~breiman/RandomForests/cc_home.htm].

- It is unexcelled in accuracy among current algorithms.
- It runs efficiently on large data bases.
- It can handle thousands of input variables without variable deletion.
- It gives estimates of what variables are important in the classification.
- It generates an internal unbiased estimate of the generalization error as the forest building progresses.
- It has an effective method for estimating missing data and maintains accuracy when a large proportion of the data are missing.
- It has methods for balancing error in class population unbalanced data sets.
- Generated forests can be saved for future use on other data.
- Prototypes are computed that give information about the relation between the variables and the classification.
- It computes proximities between pairs of cases that can be used in clustering, locating outliers, or (by scaling) give interesting views of the data.
- The capabilities of the above can be extended to unlabeled data, leading to unsupervised clustering, data views and outlier detection.
- It offers an experimental method for detecting variable interactions.

Download and Load the Data

First, download and load the data from source in to R.

```
if (!file.exists("pml-training.csv")) download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv", method = "curl")
if (!file.exists("pml-testing.csv")) download.file("https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv", method = "curl")
train <- read.csv("pml-training.csv", na.strings = c("#DIV/0!", "", "NA"))
test <- read.csv("pml-testing.csv", na.strings = c("#DIV/0!", "", "NA"))
```

Check the dimensions of two dataset

```
dim(train)
```

```
[1] 19622 160
```

```
dim(test)
```

```
[1] 20 160
```

Model Building

Split the data

In order to further cross validation, I'll split the **train** data by $p = 0.75$ using the `createDataPartition()` function in the *caret* package.

```
library(caret)
set.seed(555)
inTrain <- createDataPartition(train$classe, p = 0.75, list = FALSE)
training <- train[inTrain, ]
testing <- train[-inTrain, ]
```

Check the dimensions

```
dim(training)
```

```
[1] 14718 160
```

```
dim(testing)
```

```
[1] 4904 160
```

Preprocessing

There are many variables have near zero variance, which can cause a poor model, these can be removed using `nearZeroVar()` function, use `saveMetrics = T` to see the results.

```
# to see the results, uncomment this line nearZeroVar(training, saveMetrics = T)

# removing variables
nsv <- nearZeroVar(training)
trainingSelected <- training[, -nsv]
dim(trainingSelected)
```

```
[1] 14718 126
```

```
library(dplyr)
trainingSelected <- select(trainingSelected, -(X:num_window))
dim(trainingSelected)
```

```
[1] 14718    120
```

Some variables have many NAs inside, removing these variables.

```
# to see these variables, uncomment this line colSums(is.na(trainingSelected))

# remove these variables
trainingSelectedRmna <- trainingSelected[, colSums(is.na(trainingSelected)) == 0]
dim(trainingSelectedRmna)
```

```
[1] 14718    53
```

Fit the model

Use the function `randomForest()`, which takes data after preprocessing `data = trainingSelectedRmna` and `method = "class"`.

```
library(randomForest)
RFmodel <- randomForest(classe ~ ., data = trainingSelectedRmna, method = "class")
RFmodel
```

Call:

```
randomForest(formula = classe ~ ., data = trainingSelectedRmna,      method = "class")
      Type of random forest: classification
      Number of trees: 500
```

No. of variables tried at each split: 7

OOB estimate of error rate: 0.42%

Confusion matrix:

| | A | B | C | D | E | class.error |
|---|------|------|------|------|------|--------------|
| A | 4184 | 1 | 0 | 0 | 0 | 0.0002389486 |
| B | 11 | 2832 | 5 | 0 | 0 | 0.0056179775 |
| C | 0 | 9 | 2557 | 1 | 0 | 0.0038955980 |
| D | 0 | 0 | 26 | 2384 | 2 | 0.0116086235 |
| E | 0 | 0 | 0 | 7 | 2699 | 0.0025868441 |

Cross Validation

Use the data **testing** splitted before to perform cross validation. The process is based on the prediction on the **testing** data, and calculate the accuracy between model and testing data using `confusionMatrix()` function.

```
prediction <- predict(RFmodel, testing, type = "class")
cm <- confusionMatrix(prediction, testing$classe)
cm
```

Confusion Matrix and Statistics

| Prediction | Reference | | | | |
|------------|-----------|-----|-----|-----|-----|
| | A | B | C | D | E |
| A | 1394 | 7 | 0 | 0 | 0 |
| B | 0 | 940 | 1 | 0 | 0 |
| C | 0 | 2 | 853 | 6 | 0 |
| D | 0 | 0 | 1 | 798 | 0 |
| E | 1 | 0 | 0 | 0 | 901 |

Overall Statistics

Accuracy : 0.9963
95% CI : (0.9942, 0.9978)
No Information Rate : 0.2845
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.9954
McNemar's Test P-Value : NA

Statistics by Class:

| | Class: A | Class: B | Class: C | Class: D | Class: E |
|----------------------|----------|----------|----------|----------|----------|
| Sensitivity | 0.9993 | 0.9905 | 0.9977 | 0.9925 | 1.0000 |
| Specificity | 0.9980 | 0.9997 | 0.9980 | 0.9998 | 0.9998 |
| Pos Pred Value | 0.9950 | 0.9989 | 0.9907 | 0.9987 | 0.9989 |
| Neg Pred Value | 0.9997 | 0.9977 | 0.9995 | 0.9985 | 1.0000 |
| Prevalence | 0.2845 | 0.1935 | 0.1743 | 0.1639 | 0.1837 |
| Detection Rate | 0.2843 | 0.1917 | 0.1739 | 0.1627 | 0.1837 |
| Detection Prevalence | 0.2857 | 0.1919 | 0.1756 | 0.1629 | 0.1839 |
| Balanced Accuracy | 0.9986 | 0.9951 | 0.9978 | 0.9961 | 0.9999 |

The accuracy is really high, which is 0.9963295.

Out-of-sample error rate

The out-of-sample error rate can be calculated using the following command.

```
outOfSample <- sum(prediction != testing$classe)/length(testing$classe)
outOfSample
```

```
[1] 0.003670473
```

Prediction

Now, use the model I built to predict the original **test** data. The results is shown below in the array format for 20 data.

```
predictTest <- predict(RFmodel, test, type = "class")
predictTest
```

```
1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
B A B A A E D B A A B C B A E E A B B B
Levels: A B C D E
```

Generate Submit Answers

This is only a chunk of codes to generate the files to submit to Coursra.

```
pml_write_files = function(x) {
  n = length(x)
  for (i in 1:n) {
    filename = paste0("problem_id_", i, ".txt")
    write.table(x[i], file = filename, quote = FALSE, row.names = FALSE, col.names = FALSE)
  }
}
pml_write_files(predictTest)
```